


The Process Concept

The Process Concept

- Process – a program in execution
 - ❖ Program
 - description of how to perform an activity
 - instructions and static data values
 - Located in Disk
 - ❖ Process
 - a snapshot of a program in execution
 - memory (program instructions, static and dynamic data values)
 - CPU state (registers, PC, SP, etc)
 - operating system state (open files, accounting statistics etc)

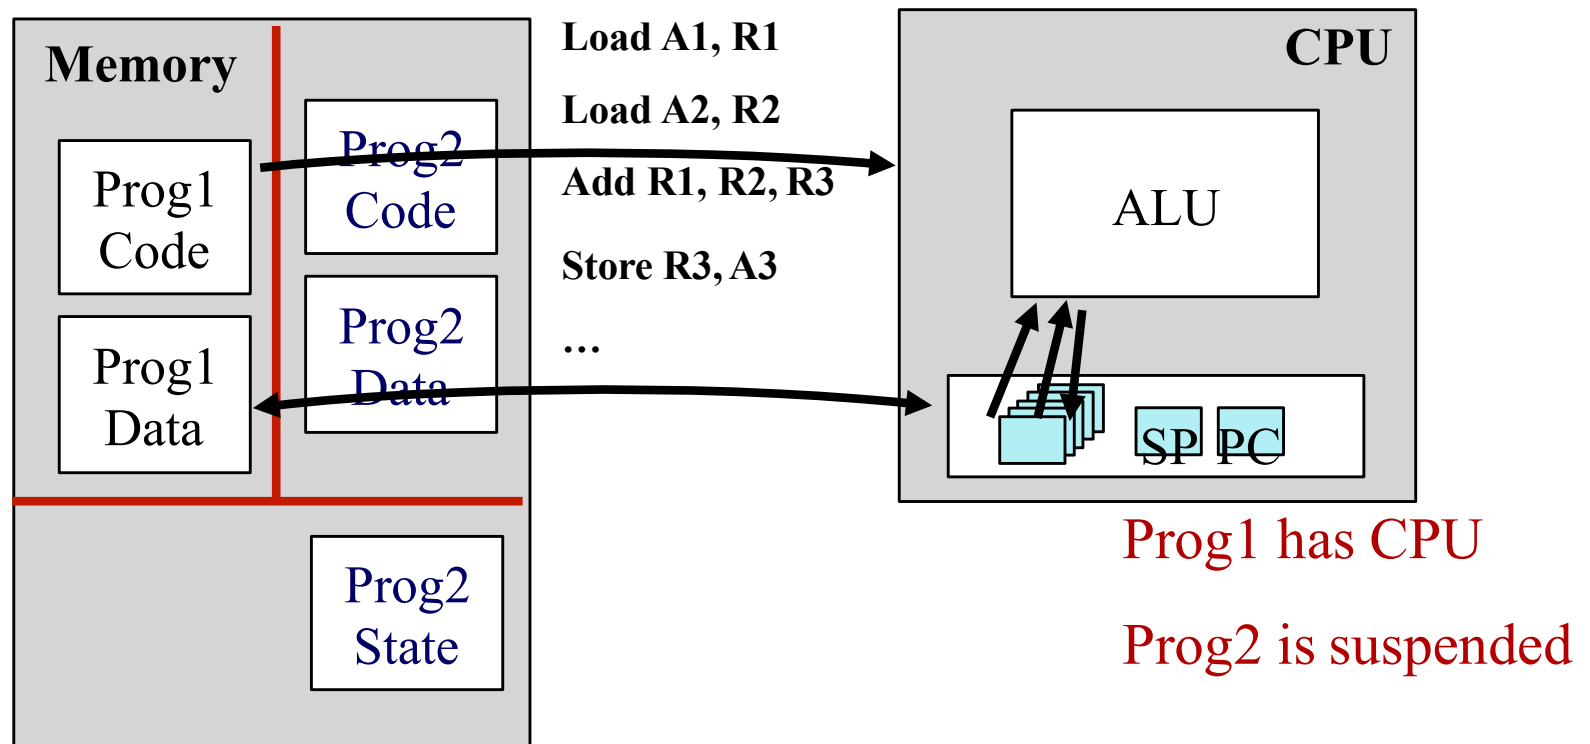
Process address space

- Each process runs in its own virtual memory *address space* that consists of:
 - ❖ *Stack space* - used for function and system calls
 - ❖ *Heap space* - used for dynamic memory allocations
 - ❖ *Data space* - Static and global variables
 - ❖ *Text* - the program code (usually read only)
- Invoking the same program multiple times results in the creation of multiple distinct address spaces



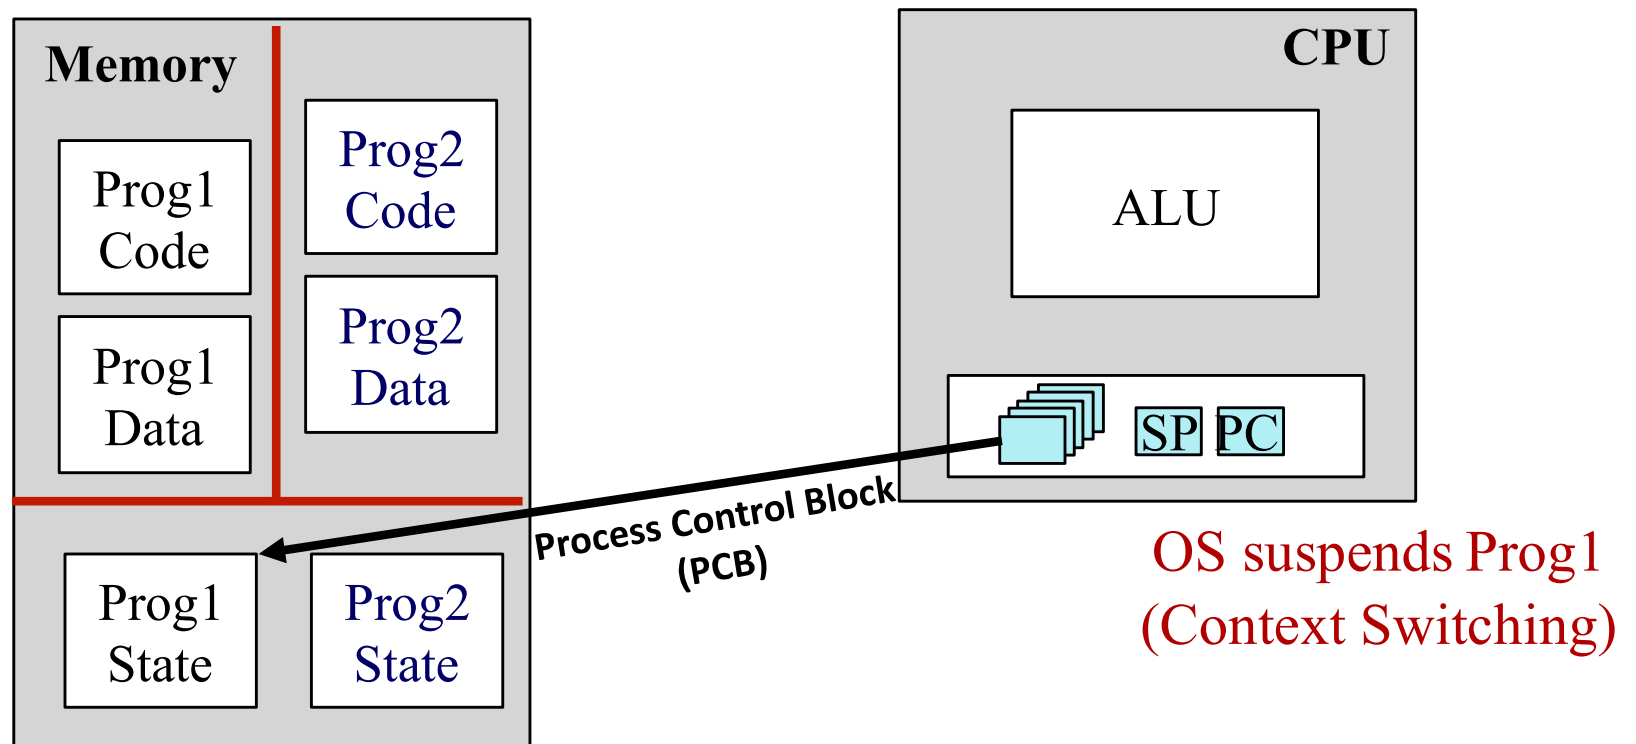
Switching among multiple processes

- Program instructions operate on operands in memory and (temporarily) in registers




Switching among multiple processes

- Saving all the information about a process allows a process to be *temporarily suspended* and later *resumed* from the same point



Switching among multiple processes

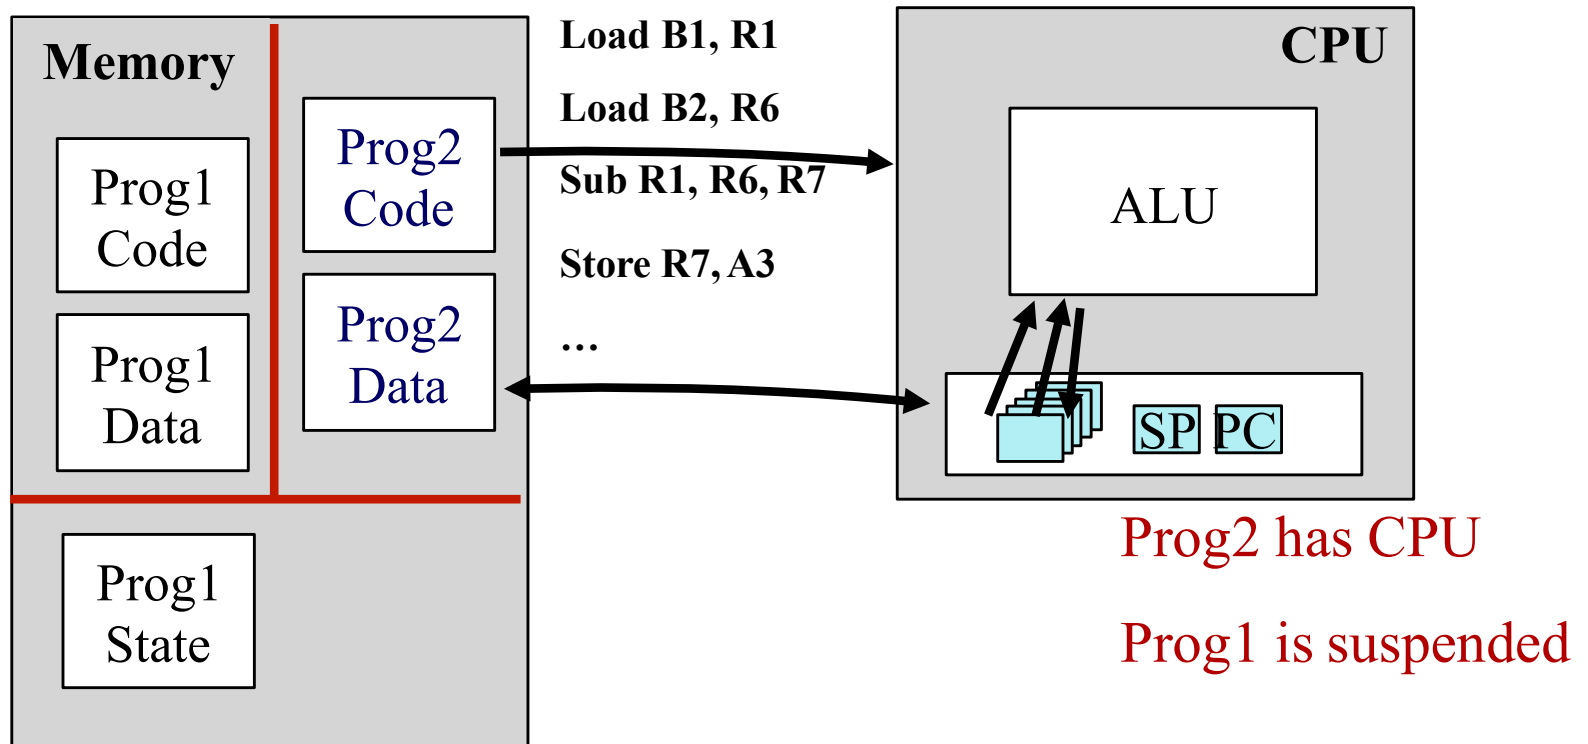
- Saving all the information about a process allows a process to be *temporarily suspended* and later *resumed*



OS resumes Prog2
(Context Switching)


Switching among multiple processes

- Program instructions operate on operands in memory and in registers




Why use the process abstraction?

One program counter




(a)

Four program counters



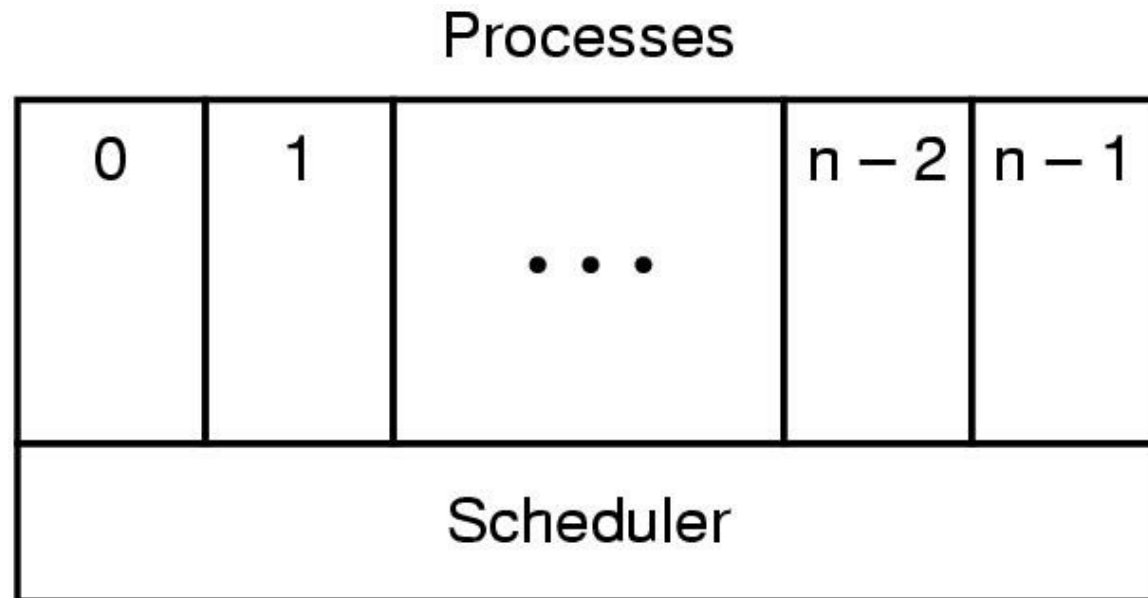
(b)



(c)

- ❑ Multiprogramming of four programs in the same address space
- ❑ Conceptual model of 4 independent, sequential processes
- ❑ Only one program active at any instant

The role of the scheduler



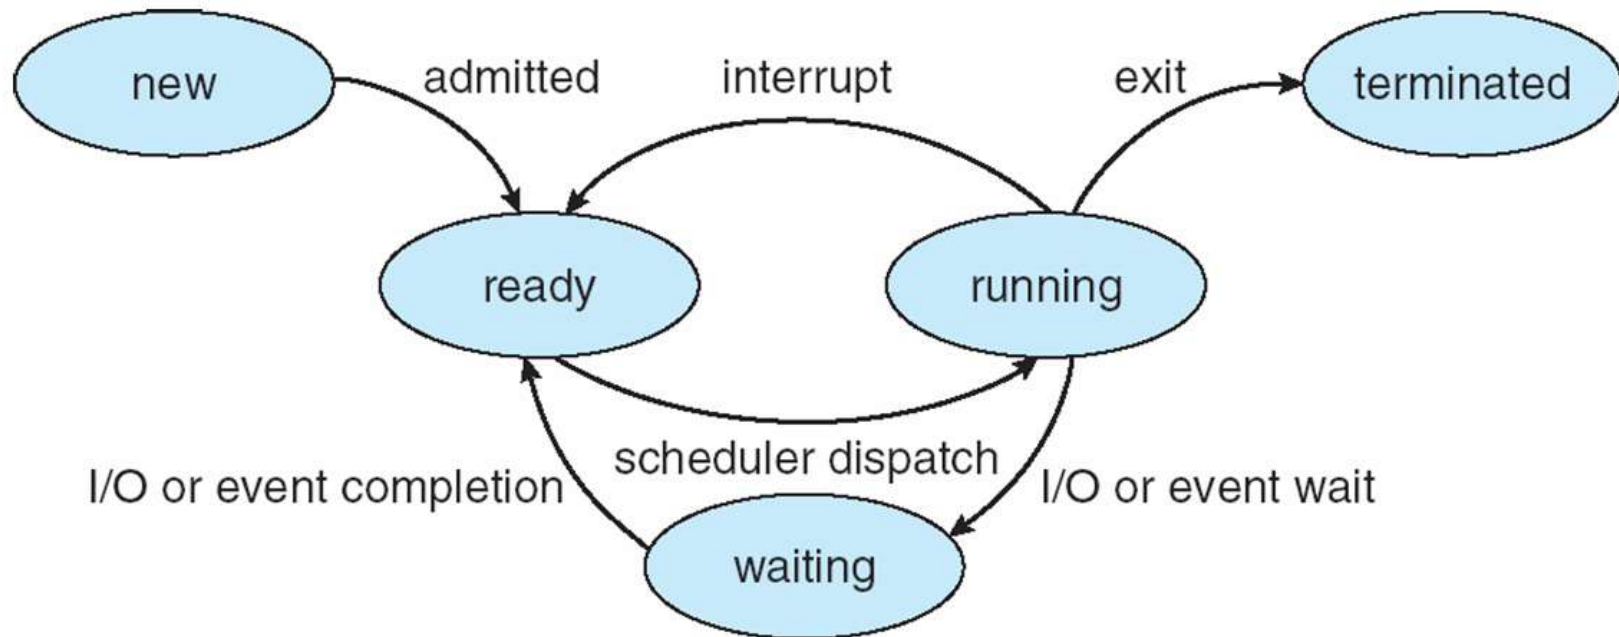
- ❑ **Lowest layer of process-structured OS**
 - ❖ handles interrupts & scheduling of processes
- ❑ **Sequential processes only exist above that layer**

Process State

As a process executes, it changes **state**

- **new**: The process is being created
- **running**: Instructions are being executed
- **waiting**: The process is waiting for some event to occur
- **ready**: The process is waiting to be assigned to a processor
- **terminated**: The process has finished execution

Diagram of Process State



How do processes get created?

Principal events that cause process creation

- ❑ System initialization
- ❑ Initiation of a batch job
- ❑ User request to create a new process
- ❑ Execution of a process creation system call from another process

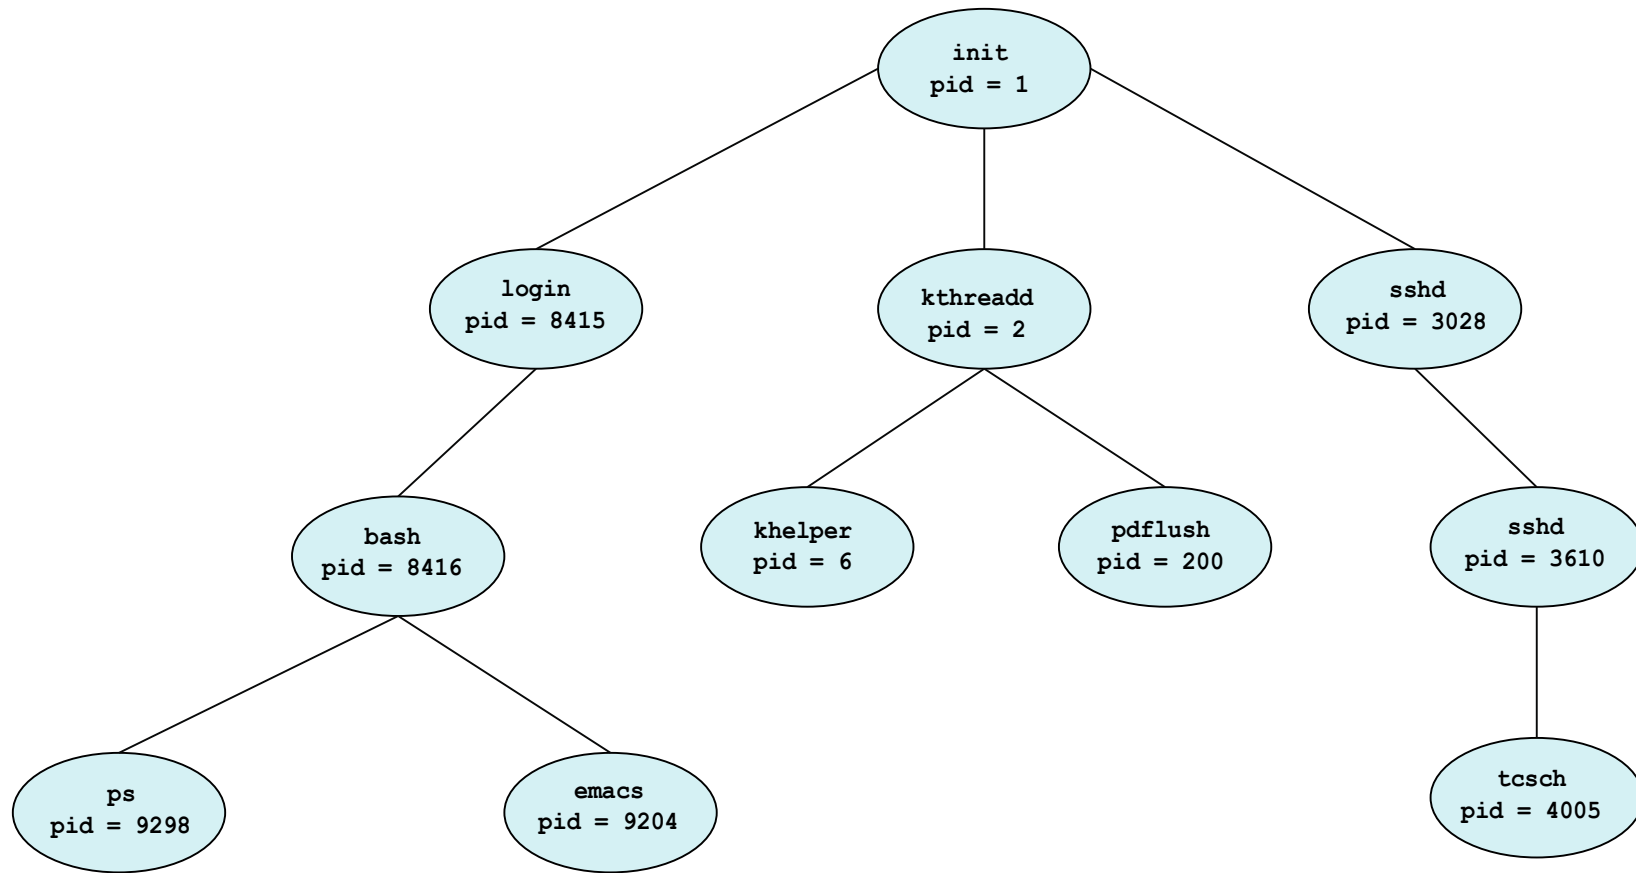
Process hierarchies

- ❑ **Parent creates a child process,**
 - ❖ special system calls for communicating with and waiting for child processes
 - ❖ each process is assigned a unique identifying number or process ID (PID)
- ❑ **Child processes can create their own child processes**
 - ❖ Forms a hierarchy
 - ❖ UNIX calls this a "process group"
 - ❖ Windows has no concept of process hierarchy

Process creation in UNIX

- ❑ All processes have a unique process id
 - ❖ *getpid()*, *getppid()* system calls allow processes to get their information
- ❑ Process creation
 - ❖ *fork()* system call creates a copy of a process and returns in both processes, but with a different return value
 - ❖ *exec()* replaces an address space with a new program
- ❑ Process termination, signaling
 - ❖ *signal()*, *kill()* system calls allow a process to be terminated or have specific signals sent to it

A Tree of Processes in Linux



Example: process creation in UNIX

csh (pid = 22)

```
...  
  
chpid = fork()  
if (chpid == 0)  
{  
    // child...  
    ...  
    exec();  
}  
else {  
    // parent  
    wait();  
... }
```


Process creation in UNIX example

csh (pid = 22)

```
...  
  
chpid = fork()  
if (chpid == 0)  
{  
    // child...  
    ...  
    exec();  
}  
else {  
    // parent  
    wait();  
... }
```

csh (pid = 24)

```
...  
  
chpid = fork()  
if (chpid == 0)  
{  
    // child...  
    ...  
    exec();  
}  
else {  
    // parent  
    wait();  
... }
```

Process creation in UNIX example

csh (pid = 22)

```
...  
  
chpid = fork()  
if (chpid == 0)  
{  
    // child...  
  
    ...  
    exec() ;  
}  
else {  
    // parent  
    wait() ;  
}
```

csh (pid = 24)

```
...  
  
chpid = fork()  
if (chpid == 0)  
{  
    // child...  
  
    ...  
    exec() ;  
}  
else {  
    // parent  
    wait() ;  
}
```

Process creation in UNIX example

csh (pid = 22)

```
...  
  
chpid = fork()  
if (chpid == 0)  
{  
    // child..  
  
    ...  
    exec() ;  
}  
else {  
    // parent  
    wait() ;  
}
```

csh (pid = 24)

```
...  
  
chpid = fork()  
if (chpid == 0)  
{  
    // child..  
  
    ...  
    exec() ;  
}  
else {  
    // parent  
    wait() ;  
}
```

Process creation in UNIX example

csh (pid = 22)

```
...  
  
chpid = fork()  
if (chpid == 0) {  
    // child...  
  
    ...  
    exec();  
}  
else {  
    // parent  
    wait();  
}  
  
...
```

ls (pid = 24)

```
//ls program  
main() {  
    //look up dir  
  
    ...  
}
```

How do processes terminate?

Conditions which terminate processes

- ❑ Normal exit (voluntary)
- ❑ Error exit (voluntary)
- ❑ Fatal error (involuntary)
- ❑ Killed by another process (involuntary)

Process Termination

- In some systems, terminating a parent process will also terminate all its child processes
- We call this case **cascading termination**
- Otherwise, if cascading termination is not enforced, a child process where its parent is terminated will be called an **Orphan process**
- Linux allows orphan processes, and assigns the **init** process to be the parent of all orphan processes
- If a child process that has terminated, but its parent hasn't yet called `wait()`, then we call it a **Zombie process**

What other process state does the OS manage?

Process management	Memory management	File management
Registers Program counter Program status word Stack pointer Process state Priority Scheduling parameters Process ID Parent process Process group Signals Time when process started CPU time used Children's CPU time Time of next alarm	Pointer to text segment Pointer to data segment Pointer to stack segment	Root directory Working directory File descriptors User ID Group ID

Example fields of a process table entry